

火焰图

一种非侵入式在线诊断技术

宦传建



高效协作 勇于担当
积极创新 用户至上

CONTENT

目录>>

01 /

火焰图是什么

02 /

有哪些优点？

03 /

如何使用？

04 /

生产案例



高效协作 勇于担当
积极创新 用户至上


生产问题

1. 原因明显

有详细的出错堆栈信息

2. 原因不明

CPU 莫名拉高？内存居高不下？服务依旧跑、也不报错



高效协作 勇于担当
积极创新 用户至上

常规方式

埋点

埋多少？侵入式修改，频繁修改上线，重启服务

摘节点

隔离生产流量，降低集群吞吐，问题复现困难



有没有一种更灵活、非侵入式的方法？

不需要修改代码、配置


不需要重启服务

不影响生产运行

.....



高效协作 勇于担当
积极创新 用户至上



基于**在线采样**生成的可视化分析工具，将时间和空间两个维度融合在一张图上，反映系统在性能方面的很多定量的统计规律。

优点

非侵入式

热插拔，随时随地，按需采集

开销很小, <3%



<http://www.brendangregg.com/flamegraphs.html>



高效协作 勇于担当
 积极创新 用户至上

如何生成？

Arthas

```
[arthas@283911]$ help
NAME      DESCRIPTION
help      Display Arthas Help
keymap    Display all the available keymap for the specified connection.
sc        Search all the classes loaded by JVM
sm        Search the method of classes loaded by JVM
classloader Show classloader info
jad       Decompile class
getstatic Show the static field of a class
monitor   Monitor method execution statistics, e.g. total/success/failure count, average rt, fail rate, etc.
stack     Display the stack trace for the specified class and method
thread    Display thread info, thread stack
trace     Trace the execution time of specified method invocation.
watch     Display the input/output parameter, return object, and thrown exception of specified method invocation
tt        Time Tunnel
jvm       Display the target JVM information
perfcounter Display the perf counter information.
ognl      Execute ognl expression.
mc        Memory compiler, compiles java files into bytecode and class files in memory.
redefine  Redefine classes. @see Instrumentation#redefineClasses(ClassDefinition...)
dashboard Overview of target jvm's thread, memory, gc, vm, tomcat info.
dump      Dump class byte array from JVM
heapdump  Heap dump
options   View and change various Arthas options
cls       Clear the screen
reset     Reset all the enhanced classes
version   Display Arthas version
session   Display current session information
sysprop   Display, and change the system properties.
sysenv    Display the system env.
vmoption  Display, and update the vm diagnostic options.
logger    Print logger info, and update the logger level
history   Display command history
cat       Concatenate and print files
echo      write arguments to the standard output
pwd       Return working directory name
mbean     Display the mbean information
grep      grep command for pipes.
tee       tee command for pipes.
profiler  Async Profiler. https://github.com/jvm-profiling-tools/async-profiler
stop      Stop/Shutdown Arthas server and exit the console.
```

```
[arthas@283911]$ profiler start
Started [cpu] profiling
[arthas@283911]$ profiler status
[perf] profiling is running for 2 seconds
[arthas@283911]$ profiler getSamples
1480
[arthas@283911]$ profiler stop
OK
profiler output file: /home/presto/data232/arthas-output/20200917-110107.svg
[arthas@283911]$ █
```



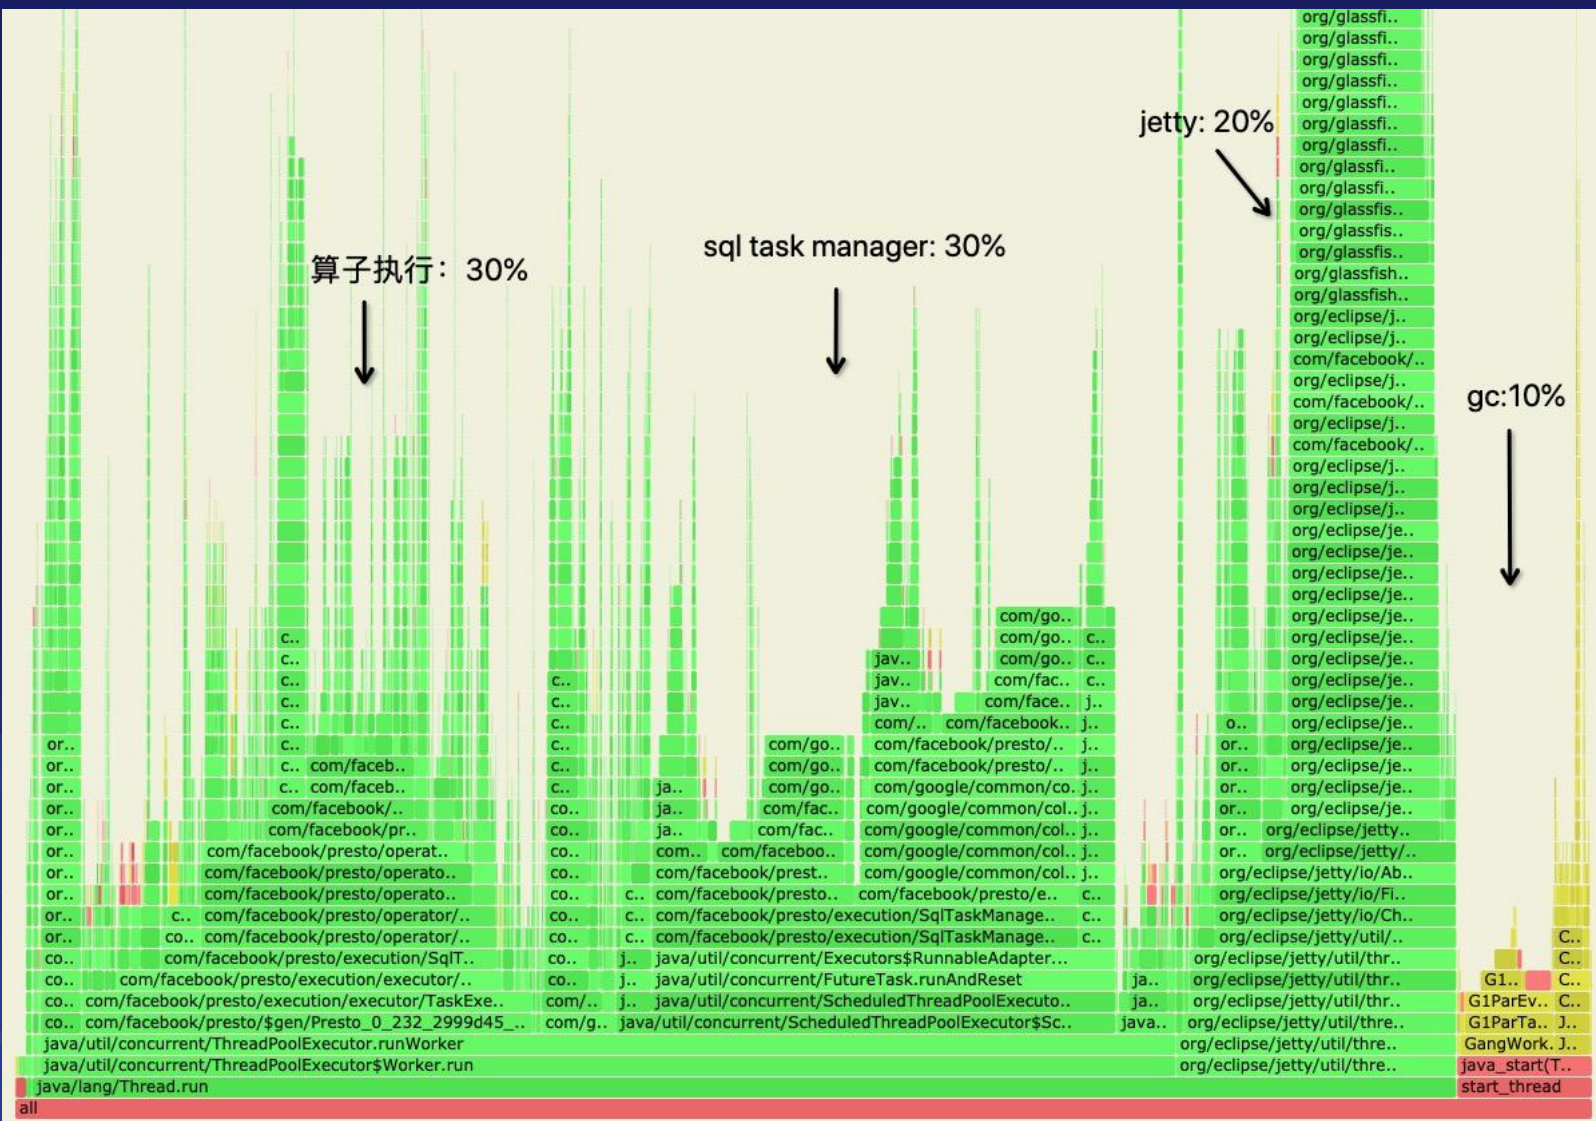
高效协作

勇于担当

积极创新

用户至上

如何看图？



x轴表示 stack 采样数量，越宽的表示出现的频率越高

y轴表示 stack 的调用深度

颜色的深浅不重要，是随机表示的

宽平顶通常意味着潜在的热点



高效协作

积极创新


勇于担当

用户至上

task manager 热点消除

task info 过期时间从 15分钟改为 5分钟

http 请求优化



worker 的 http 请求主要用于 task 之间交换数据

集群不宜过大，主节点会为每个 task 创建两个轮询任务：轮询 task 执行状态、定时更新 task info

小查询居多，减少不必要的 task

buffer 不宜过大，启发式公式会根据 buffer 的剩余容量发起更多的 http 请求





高效协作

勇于担当

积极创新

用户至上

airlift http request stats - OneMinute.Rate




高效协作 勇于担当
积极创新 用户至上

<https://github.com/jvm-profiling-tools/async-profiler>



高效协作 勇于担当
积极创新 用户至上

An isometric illustration of a ZTO Express warehouse and sorting facility. The main building is blue with the ZTO logo and the slogan '成为全球一流的综合物流服务商' (Become a first-class comprehensive logistics service provider globally) on its facade. In the foreground, a conveyor belt system is shown with several brown cardboard boxes being sorted. A worker in a blue uniform is standing near the conveyor. The entire scene is set against a dark blue background.



高效协作 勇于担当
 积极创新 用户至上

```
public static void copy(Configuration from, Configuration to) from: "Configuration: " to: "Configuration: "
{
    for (Map.Entry<String, String> entry : from) { from: "Configuration: "
        to.set(entry.getKey(), entry.getValue());
    }
}

ConfigurationUtils > copy()

is
loadDefaults = false
updatingResource = {ConcurrentHashMap@20016} size = 438
properties = {Properties@20017} size = 438
  "io.bytes.per.checksum" -> "512"
  "fs.s3a.connection.timeout" -> "50000"
  "hadoop.security.kms.client.encrypted.key.cache.expiry" -> "43200000"
  "dfs.image.transfer.bandwidthPerSec" -> "0"
  "hadoop.security.sensitive-config-keys" -> "password$,fs.s3.*[Ss]ecret.?[Kk]ey,fs.azure.account.key.*,dfs.webhdfs.oauth2.[a-z].+token,hadoop.security"
  "dfs.block.access.token.lifetime" -> "600"
  "dfs.namenode.resource.du.reserved" -> "104857600"
  "hadoop.proxyuser.hive.hosts" -> ""
  "fs.AbstractFileSystem.file.impl" -> "org.apache.hadoop.fs.local.LocalFs"
  "fs.du.interval" -> "600000"
  "fs.AbstractFileSystem.har.impl" -> "org.apache.hadoop.fs.HarFs"
  "dfs.domain.socket.path" -> "/var/run/hdfs-sockets/dn"
  "dfs.balancer.max-no-move-interval" -> "60000"
  "hadoop.ssl.keystores.factory.class" -> "org.apache.hadoop.security.ssl.FileBasedKeyStoresFactory"
  "fs.s3a.multipart.size" -> "104857600"
  "hadoopn http authentication kerberos keytab" -> "${user.home}/hadoopn keytab"
```

```
public void set(String name, String value, String source) {
    Preconditions.checkArgument(
        expression: name != null,
        errorMessage: "Property name must not be null");
    Preconditions.checkArgument(
        expression: value != null,
        errorMessage: "The value of property " + name + " must not be null");
    name = name.trim();
    DeprecationContext deprecations = deprecationContext.get();
    if (deprecations.getDeprecatedKeyMap().isEmpty()) {
        getProps();
    }
    getOverlay().setProperty(name, value);
    getProps().setProperty(name, value);
    String newSource = (source == null ? "programmatically" : source);



    if (!isDeprecated(name)) {
        updatingResource.put(name, new String[] {newSource});
        String[] altNames = getAlternativeNames(name);
        if (altNames != null) {
            for (String n: altNames) {
                if (!n.equals(name)) {
                    getOverlay().setProperty(n, value);
                    getProps().setProperty(n, value);
                    updatingResource.put(n, new String[] {newSource});
                }
            }
        }
    }
}
```



Presto 调度阶段访问 HDFS



高效协作 勇于担当
积极创新 用户至上



* off-CPU 维度的火焰图：进程/线程不运行在任何 CPU 上的时间，锁、竞争、I/O 等待。。。



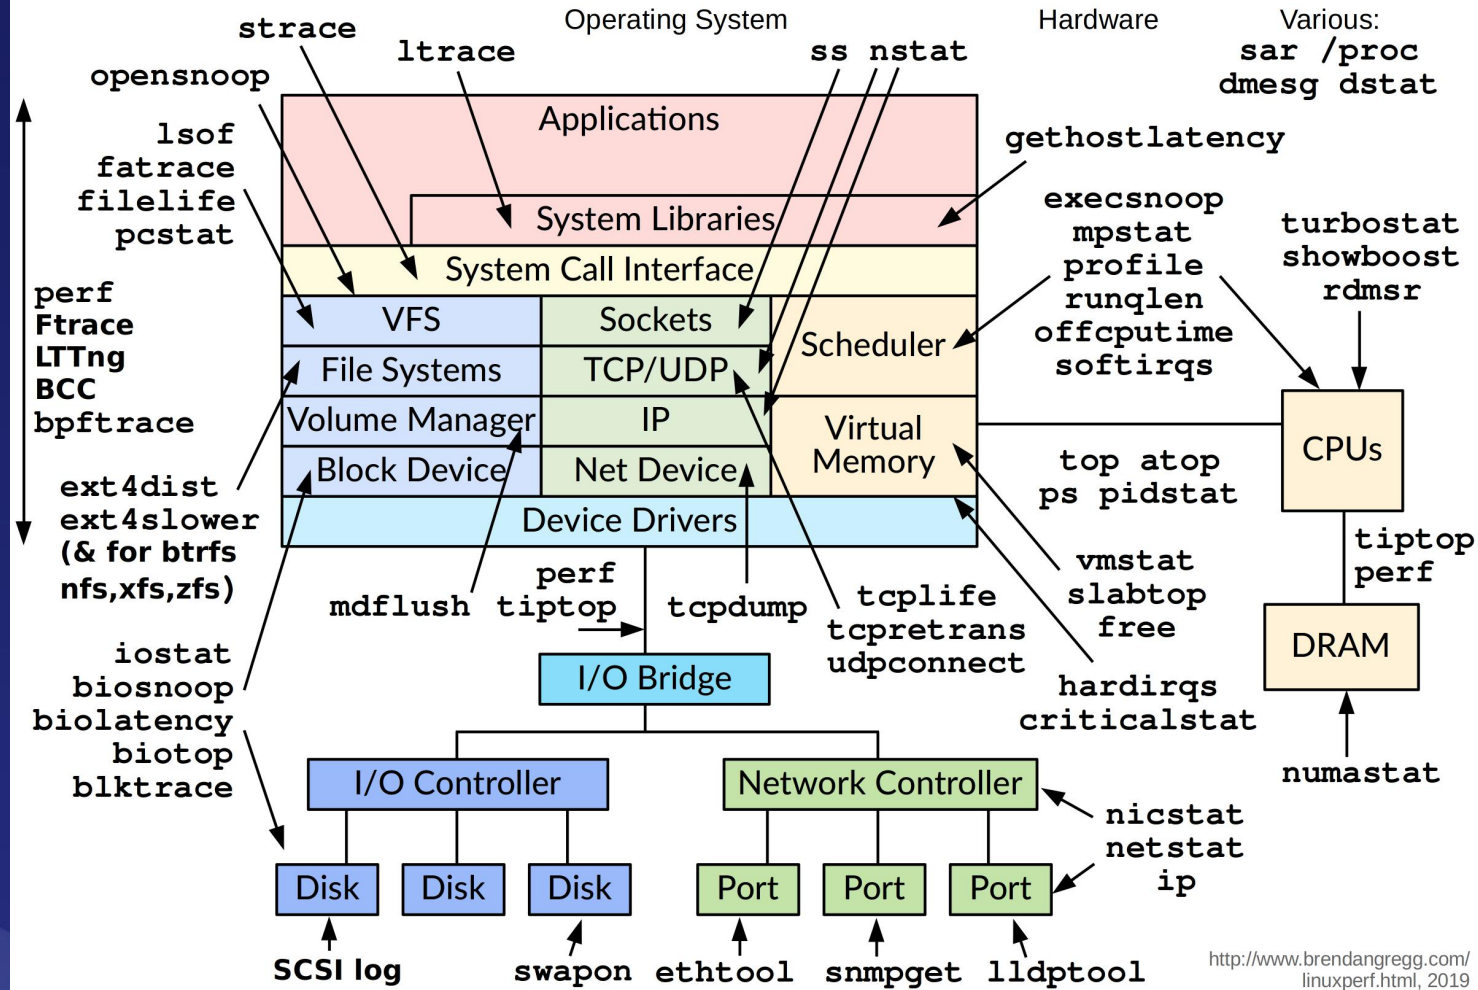
高效协作 勇于担当
积极创新 用户至上

动态追踪 or 静态分析？



高效协作 勇于担当
积极创新 用户至上

Linux Performance Observability Tools



高效协作 勇于担当
积极创新 用户至上



有什么想知道的？

What would You like to know?

我想知道...

I wank to know...



高效协作
积极创新

勇于担当
用户至上



谢谢

THANK YOU

 中通科技